TestOrCode: Mid-Submission Progress Report

Intent-grounded triage of CI test failures with scoped change memory and test-integrity checking. Team Agent Zero. Mid-submission, 30 October 2026. All numbers are preliminary and come from saved runs.

1. Executive summary and current status

TestOrCode decides, for each failing pytest test after a code change, whether to UPDATE_TEST, FIX_CODE or ESCALATE, and checks that the repair does not weaken the test suite. The submitted proposal describes four logical agents (Triage, Intent Analyst, Repairer, Integrity Guard), intent-grounded retrieval, persistent memory backed by executable probes, and a mutation-based integrity check.

- **Working end to end:** the whole pipeline (Triage, Intent Analyst with BM25 retrieval, Repairer with a restricted patch tool, Integrity Guard, probe-backed memory, simulated maintainer) runs from a CLI against real boltons episodes with a hosted model.
- **Conditions implemented:** FULL, B1 (single tool-using agent with RAG, no memory), A1 (no intent retrieval), A2 (no memory), A3 (no Integrity Guard), plus the always-UPDATE_TEST and always-FIX_CODE sanity baselines. All use one episode format and one result format.
- **Pilot dataset:** 13 episodes on one repository (9 development, 4 frozen test), built from real commits.
- **Real runs:** all conditions were run live on the 9 development episodes and, once, on the 4 frozen test episodes (one seed, one model). Test-split results are in Section 6.
- **Not claimed:** no statistical result, no conclusion that any condition is better than another. The pilot is too small; it shows the experiment runs and exposes the failure modes the final 15 days must address.

2. Research question and proposed architecture

Research question. On chronological streams of CI failures in small Python libraries, does a role-separated pipeline that (i) grounds each decision in retrieved change-intent evidence, (ii) remembers approved behaviour changes as scoped records with executable probes, and (iii) checks test integrity after repair make fewer wrong-side repairs and weaken tests less than a single tool-using agent with the same tools, model and budget?

Architecture. One Python orchestrator (a deterministic state machine) with four role-specific prompts. Triage runs the suite, confirms stability and groups failures by function. The Intent Analyst retrieves evidence (BM25 over commit messages, docs, changelog, docstrings and source comments, plus probe-gated memory) and states a hypothesis with cited evidence ids. The decision policy requires cited positive evidence for UPDATE_TEST and cited expected-behaviour evidence for FIX_CODE, otherwise ESCALATE. The Repairer applies a scoped patch through a restricted patch tool. The Integrity Guard re-runs tests, mutants, weakening checks and a semantic review, and is the only role that writes memory, after simulated maintainer approval. All conditions go through one model client and share budgets.

Safety design. Path and scope checks, the allowed-file policy, and an expected-literal-only rule for test patches are enforced by the host-side tools, not by the prompt. Tests run in a bubblewrap sandbox without network. Gold labels never reach the agent: the harness keeps them, and agent-view files contain none by construction.

3. Implemented components versus planned components

Component	Status	Evidence / note
Four roles and orchestrator	Implemented	Live runs on 9 dev and 4 test episodes, FULL/A1/A2/A3
B1 single-agent baseline	Implemented	Same tools, standard RAG, no memory; live runs
BM25 retrieval, gold-evidence Recall@k	Implemented	docs/retrieval.md; hybrid embeddings planned
Probe-backed scoped memory	Implemented	Probes run on current code before reuse; stale records flagged
Restricted patch tool, weakening checks	Implemented	Tested on real episode states

Component	Status	Evidence / note
Mutation-based Guard check	Implemented, coarse	Cached mutants; small functions give 1-2 mutants (dev-06: none killed)
A1, A2, A3 switches	Implemented, run on dev and test	Single seed; A1 escalated every decision on dev and test
Experiment runner, metrics, report tables	Implemented	Resume, retry-failed, post-hoc integrity audit
Docker sandbox	Implemented, untested	All runs used bubblewrap; Docker needs daemon access
Tool-fault injection (scenario T)	Planned	Not wired; fault plans exist in two dev episodes
git_log / git_show / git_blame tools	Stubs	Not needed by current episodes; dev-01 stays provisional
Embeddings and reranking	Planned	Optional; BM25 first, as in the 5-day plan
72 episodes, 3 held-out libraries, 2 seeds	Planned	Deferred to the 15-day plan
Statistics (bootstrap, McNemar)	Planned	Not meaningful at pilot size
Fine-tuning, frontend	Dropped	Outside scope by the plan

4. Dataset: sources, construction, labels and actual size

The pilot uses boltons (BSD-3-Clause), pinned at commit d31d821d. It has 9 development and 4 frozen test episodes, kept in separate chronological streams with disjoint commits, tests, functions and modules. This is a pilot of one library, not the 72-episode, three-library set of the proposal. Outdated-test episodes (S1, S3) come from real commits that changed source and tests; regression episodes (S2, S4) inject a mutant into real code; S5 strips the intent evidence. Every episode was verified by running the code: the CI state fails the same tests in 5 of 5 runs and, where a gold repair exists, it turns the suite green. The S5 episode (dev-01) is provisional.

dev split episodes

Episode	Scenario	Difficulty	Failing tests	Expected action(s)
dev-01	S5	medium	2	ESCALATE
dev-02	S2	easy	1	FIX_CODE
dev-03	S2	medium	2	FIX_CODE
dev-04	S2	easy	1	FIX_CODE
dev-05	S1	easy	1	UPDATE_TEST
dev-06	S1	easy	1	UPDATE_TEST
dev-07	S4	hard	3	FIX_CODE
dev-08	S3	hard	4	FIX_CODE, UPDATE_TEST
dev-09	S1	easy	1	UPDATE_TEST

dev split, per failing test: gold direction FIX_CODE 10, UPDATE_TEST 6; expected action ESCALATE 2, FIX_CODE 10, UPDATE_TEST 4.

test split episodes

Episode	Scenario	Difficulty	Failing tests	Expected action(s)
test-01	S2	easy	1	FIX_CODE
test-02	S2	easy	1	FIX_CODE
test-03	S3	hard	3	FIX_CODE, UPDATE_TEST
test-04	S4	hard	1	FIX_CODE

test split, per failing test: gold direction FIX_CODE 4, UPDATE_TEST 2; expected action FIX_CODE 4, UPDATE_TEST 2.

The development split is FIX_CODE-heavy. Gold evidence ids exist for 5 of the 13 episodes (all UPDATE_TEST); recall@k is undefined elsewhere and is shown as '-'.

5. B1 baseline, ablations and experiments completed

Condition	Definition	Status
FULL	All four roles, intent retrieval, probe-gated memory, Guard	Run live (dev, test)
B1	One tool-using agent, standard RAG, no persistent memory	Run live (dev, test)
A1	FULL without intent retrieval: no commit message, corpus or search_repo; memory kept	Run live (dev, test)
A2	FULL without persistent memory	Run live (dev, test)
A3	FULL without the independent Integrity Guard (test-level verification only)	Run live (dev, test)
always UPDATE_TEST / FIX_CODE	Deterministic sanity baselines, no model	Run (dev, test)

The proposal's A3 (no Integrity Guard) is used, not the detailed idea's coarse-memory ablation. Tool-level refusal of weakening edits applies in every condition, so A3's weakening rate is 0 by construction. B0, B2, A4 and A5 of the detailed idea are not implemented.

6. Preliminary quantitative results and limitations

Development split, 9 episodes, seeds [0], model gemini-3.5-flash-lite. Preliminary.

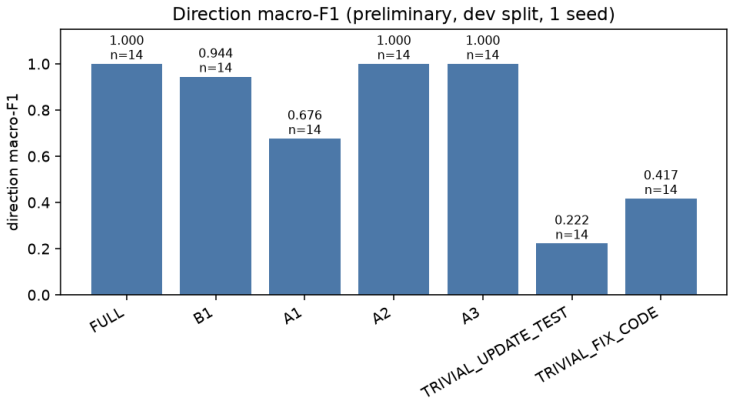
Condition	Episodes run	Direction macro-F1	Direction acc.	Final-action acc.	Correct-repair rate	Masked regressions
FULL	9	1.000	14/14 (1.00)	0.562	3/6 (0.50)	0/10 (0.00)
B1	9	0.944	12/14 (0.86)	0.250	0/4 (0.00) +2 unverif.	0/10 (0.00)
A1	9	0.676	11/14 (0.79)	0.125	-	0/10 (0.00)
A2	9	1.000	14/14 (1.00)	0.688	5/7 (0.71)	0/10 (0.00)
A3	9	1.000	14/14 (1.00)	0.438	0/3 (0.00) +3 unverif.	0/10 (0.00)
always UPDATE_TEST	9	0.222	4/14 (0.29)	0.250	-	10/10 (1.00)
always FIX_CODE	9	0.417	10/14 (0.71)	0.625	-	0/10 (0.00)

- Integrity audit, A1: 0 accepted repair(s) in 9 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, A2: 6 accepted repair(s) in 9 run(s); weakened 0/2 measured test patches; 0 accepted on the wrong side; 1 not fully measured; mutant-kill retention 4/6 reference-killed mutants over 1 test patch(es)
- Integrity audit, A3: 4 accepted repair(s) in 9 run(s); weakened 0/2 measured test patches; 0 accepted on the wrong side; 1 not fully measured; mutant-kill retention 4/6 reference-killed mutants over 1 test patch(es)
- Integrity audit, B1: 4 accepted repair(s) in 9 run(s); weakened 0/3 measured test patches; 0 accepted on the wrong side; 1 not fully measured; mutant-kill retention 7/9 reference-killed mutants over 2 test patch(es)
- Integrity audit, FULL: 4 accepted repair(s) in 9 run(s); weakened 0/1 measured test patches; 0 accepted on the wrong side; 1 not fully measured; mutant-kill retention not measurable
- Integrity audit, TRIVIAL_FIX_CODE: 0 accepted repair(s) in 9 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, TRIVIAL_UPDATE_TEST: 0 accepted repair(s) in 9 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable

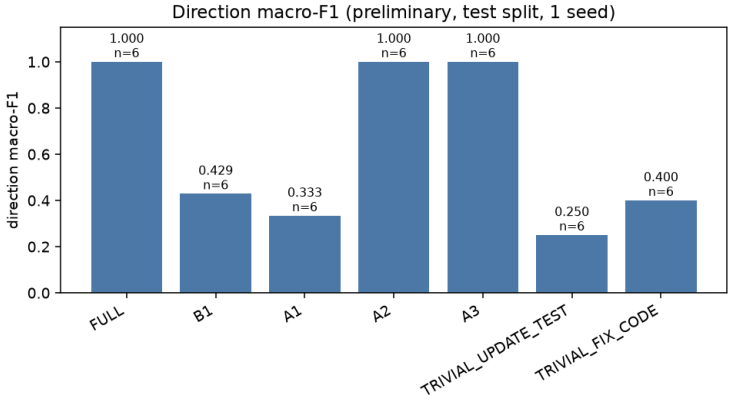
Frozen test split, 4 episodes, seeds [0], model gemini-3.5-flash-lite. Preliminary.

Condition	Episodes run	Direction macro-F1	Direction acc.	Final-action acc.	Correct-repair rate	Masked regressions
FULL	4	1.000	6/6 (1.00)	0.167	1/2 (0.50)	0/4 (0.00)
B1	4	0.429	3/6 (0.50)	0.167	0/2 (0.00) +1 unverif.	0/4 (0.00)
A1	4	0.333	3/6 (0.50)	0.000	-	0/4 (0.00)
A2	4	1.000	6/6 (1.00)	0.333	1/2 (0.50)	0/4 (0.00)
A3	4	1.000	6/6 (1.00)	0.167	0/1 (0.00) +1 unverif.	0/4 (0.00)
always UPDATE_TEST	4	0.250	2/6 (0.33)	0.333	-	4/4 (1.00)
always FIX_CODE	4	0.400	4/6 (0.67)	0.667	-	0/4 (0.00)

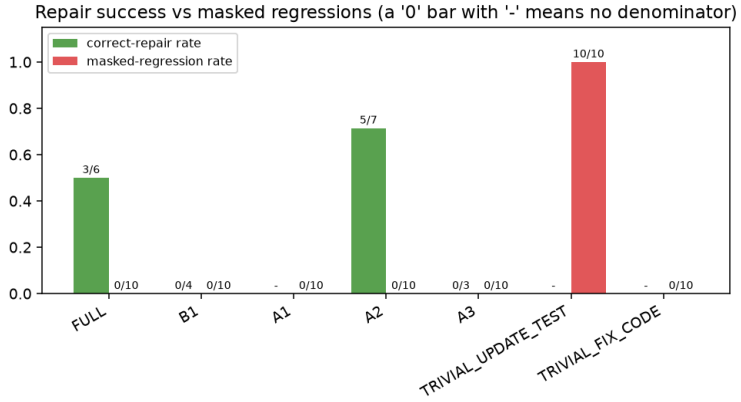
- Integrity audit, A1: 0 accepted repair(s) in 4 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, A2: 2 accepted repair(s) in 4 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, A3: 1 accepted repair(s) in 4 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, B1: 1 accepted repair(s) in 4 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, FULL: 1 accepted repair(s) in 4 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, TRIVIAL_FIX_CODE: 0 accepted repair(s) in 4 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable
- Integrity audit, TRIVIAL_UPDATE_TEST: 0 accepted repair(s) in 4 run(s); weakened 0/0 measured test patches; 0 accepted on the wrong side; 0 not fully measured; mutant-kill retention not measurable



Direction macro-F1 on the development split



Direction macro-F1 on the frozen test split



Repair success against masked regressions

How to read the results.

- Direction is scored from the agent's hypothesis, so it is generous. An escalation counts by the side it blamed. Final-action accuracy is lower because conditions escalate often; read it with direction.
- Correct-repair rates have denominators of 1 to 7 repairs. 'Unverifiable' means a repair was accepted but no trusted probe confirmed it (A3 and B1 have no Guard).
- Masked regressions are 0 for every model condition because none edited a test to hide a regression. The metric only separates conditions from always-UPDATE_TEST. It does not show that the Guard is needed.
- A1 escalated every decision, so its repair rate is undefined; its direction score comes from the hypothesis alone.
- The post-hoc audit found no weakened and no wrong-side accepted repair in any condition. Mutant-kill retention is pooled over only the test patches whose reference tests kill a mutant (raw counts above).
- **Limitations:** one seed, one model, 9 (+4) episodes, one library; no confidence intervals; citation support, stale-evidence rate and memory over-generalisation are not computed; Docker untested; scenario T not wired.

One success and one failure from the logs. dev-03 (FULL): the Analyst cited the commit message and the docstring, the source patch applied, the Guard's mutants passed, the suite was green and a memory record was written. dev-05 (FULL, A2, A3): the hypothesis was correct (outdated test) but both repair attempts edited more than expected literals and the patch-content rule refused them, so the run escalated although the direction was right. This is the strict safety rule working, and its cost.

Reading the frozen test-split runs (hand review of `review_runs.py` output on `results/day4/pilot-test-5cond`; one run per cell, so these are observations, not effects).

- **Success, test-01 (S2, gold FIX_CODE), FULL:** the Analyst cited the commit message, the Repairer's source patch applied, the Guard approved after the mutants were run, the suite passed and a memory record was written. A2 and A3 reached the same repair.
- **Failure, test-02 (S2, gold FIX_CODE), all five conditions:** every condition escalated. Each hypothesised a regression correctly but cited no expected-behaviour evidence, so the decision policy escalated. Final-action accuracy is therefore 0 on this episode for all of them, although direction was right.
- **Failure, test-03 (S3 mixed):** for the outdated half, FULL and A2 had their test patches refused by the literal-only rule (the model changed more than expected values), while A3 and B1 produced diffs that did not apply. All escalated. For the regression half only A2 produced a Guard-approved source fix; FULL and A3 escalated for missing evidence.
- **Decoy, test-04 (S4, gold FIX_CODE):** FULL and A3 read the decoy memory record (its probe passed) but hypothesised a regression without citing it, and escalated. A1 cited the decoy record for an outdated-test hypothesis (the wrong direction): the only run in which a hypothesis followed the decoy, and its escalation meant no wrong-side edit was made. B1, with no memory, fixed the code correctly.

- **B1 on test-01:** it hypothesised an outdated test for a real regression. Its two test edits were refused by the patch rule, so the masked regression was blocked by a tool rule, not by B1's judgment.
- **What this means:** on 4 episodes the ordering by direction macro-F1 (FULL, A2, A3 above B1 and A1) is driven by a few tests: B1 is wrong on test-01 and on both test-03 rotate_file tests, A1 on both rotate_file tests and on test-04, and final-action accuracy is low everywhere because the evidence and patch-content rules turn most decisions into escalations. This is evidence about the policy's conservatism, not about which design is better.

7. Challenges faced and mitigation plans

Challenge	Mitigation
Free-tier quota and 503 overload (7 of 19 Day-2 runs were api_failure)	Key pool with quota memory, retries, api_failure never scored, --retry-failed with memory-order guard
Prompt contamination: Day-2 repair prompts contained dev-05 content	Fixed in repair-v5; a test fails if a prompt shares 20 characters with any pilot episode; old records kept with a notice
A1 side channel: commit message reachable by retrieval or read_file	Corpus, search_repo and non-code read_file withheld in A1; checked by tests
Strict test-patch rule causes escalations (dev-05)	Report as a cost of safety; consider more allowed setup edits only on the development split
Mutation strength unmeasurable for small functions (dev-06: reference test kills 0 of 2 mutants)	More mutant operators and a separate 'strength unmeasurable' class in the final 15 days
A1 and probe-less repairs are scored ambiguously	Report direction and final-action accuracy side by side; decide the scoring rule before the final run
Bubblewrap only; Docker daemon not accessible	Test Docker on a machine with access, keep bwrap as the fallback
Small pilot cannot support statistics	Scale to three libraries and 72 episodes, two seeds, paired bootstrap

8. Updated repository and reproducibility

- Repository: this project's `main` branch; `README.md` has the setup, `docs/interfaces.md` the interfaces and `docs/day*-*.md` the per-day logs.
- Setup: Python 3.12 virtual environment, `pip install -e .[dev,report]`, `.env` from `.env.example` (keys are never committed), a clone of boltons at the pin and its pilot venv, bubblewrap.
- Offline checks: `python -m pytest -q` (no API calls). Live runs: `scripts/run_experiment.py`, then `scripts/audit_integrity.py`, `scripts/summarize_results.py`, `scripts/make_report.py`.
- Every table in this report is produced by `docs/report/build_report.py` from `results/**/runs.jsonl`.

9. Completion plan, 31 October to 14 November

Dates	Work
Oct 31 - Nov 2	Stabilise dataset and environments; episode-construction scripts; select the three held-out libraries; validate gold labels; pin dependencies; development stream
Nov 3 - 5	Finish core features: hybrid retrieval, memory, Guard, fault recovery; fix the A1 and escalation scoring; freeze model and configuration on development data
Nov 6 - 9	Full evaluation: 72-episode test set for FULL and B1 (two seeds), then A1, A2, A3 and sanity baselines, checkpointed. Throughput check first: about 74 s per run and 400k tokens for 45 pilot runs
Nov 10 - 11	Metrics, paired statistics where feasible, failure analysis, test-weakening and incorrect-repair review
Nov 12 - 13	Final report, plots, reproducibility package, demo
Nov 14	Clean setup test, repository check, backups, final release; Nov 15 is submission only

10. Demo and known limitations

Demo: `bash scripts/demo.sh` runs offline, with no API call, and writes nothing under `results/`. It rebuilds the CI state of dev-03 and shows its real failing tests, then replays the saved FULL run (evidence cited, decision, accepted patch, Guard approval), the saved B1 and A3 runs on the same episode, a no-model baseline through `scripts/run_episode.py`, and the results table rebuilt from the saved records. The commands, expected output, an actual transcript and a terminal recording (`scriptreplay`) are in `docs/demo/`. A live run with the model is possible (`scripts/run_episode.py --condition A2 --episode dev-03`) but spends API quota and was not used for the demo. Known limitations are listed in Sections 3 and 6.